

# Paradigms in Distributed Computing

## Grado en Ingeniería Informática

Pablo García Sánchez

Departamento de Ingeniería Informática

Universidad de Cádiz

Pablo García Sánchez

*Based on the work by Guadalupe Ortiz, Mei Ling-Liu, Marteen Van Steem and A. Tanenbaum*



Curso 2017 – 2018

# Index

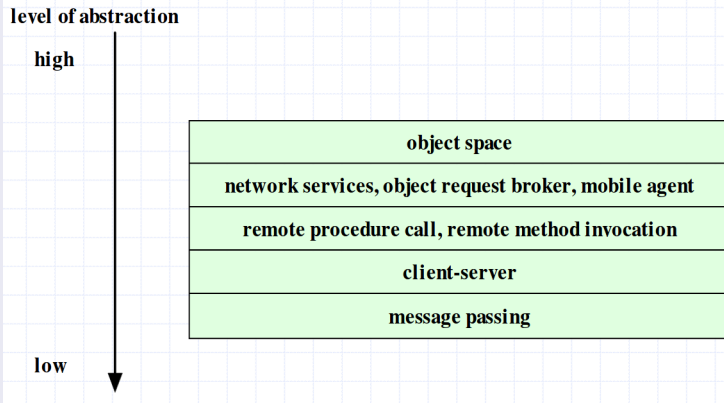
- 1 Introduction
- 2 Message Passing Paradigm
- 3 Client-Server paradigm
- 4 Peer-to-peer
- 5 Message System Paradigm
- 6 Remote Procedure Call
- 7 Distributed Objects

# Section 1 | Introduction

# Paradigm

- Definition of **paradigm**: a pattern, example or model.
- Characteristics of distributed applications vs. conventional applications (one machine)
  - Interprocess Communication
  - Event Synchronization
- **Abstraction** in engineering is realized with the provision of tools or facilities which allow software to be built without the developer having to be cognizant of some of the underlying complexities.

# Distributed Application Paradigms



## Section 2 | Message Passing Paradigm


# Message Passing Paradigm


- The most fundamental paradigm for distributed applications.
- A process sends a message representing a request.
- The message is delivered to a receiver, which processes the request, and sends a message in response.
- The reply may trigger a further request, which leads to a subsequent reply, and so forth.

# Message Passing

**Process A**

**Process B**



 a message

**Message passing**



# Message Passing Paradigm

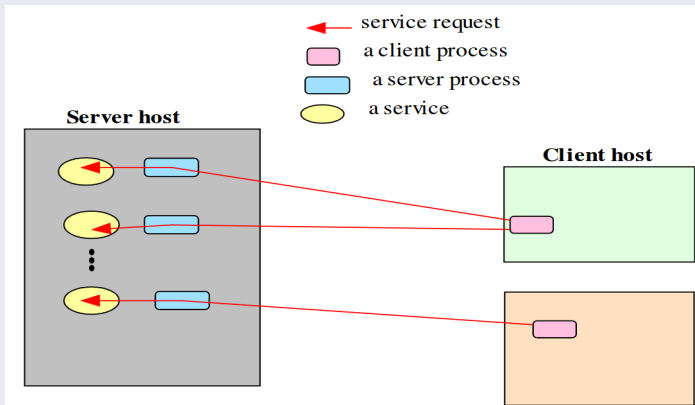
- Basic operations required to support the basic message passing paradigm are **send**, and **receive**.
- For **connection-oriented** communication, the operations **connect** and **disconnect** are also required.
- With the abstraction provided by this model, the interconnected processes perform input and output to each other, in a manner similar to file I/O. The I/O operations encapsulate the detail of network communication at the operating-system level.
- The socket application programming interface is based on this paradigm, as we know :)

## Section 3 | Client-Server paradigm

# Introduction to Client-Server

- The most prevalent model for distributed computing protocols.
- It is the basis of all distributed computing paradigms at a higher level of abstraction.
- It is [service-oriented](#), and employs a request-response protocol.

# Client-Server



# Client-Server

- A server process, running on a server host, provides access to a service.
- A client process, running on a client host, accesses the service via the server process.
- The interaction of the process proceeds according to a [protocol](#).

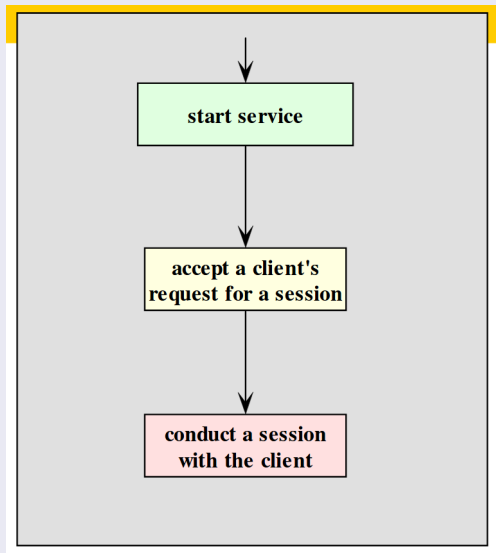
# Client-Server: applications and services

- An application based on the client-server paradigm is a client-server application.
- On the Internet, many services are Client-server applications. These services are often known by the protocol that the application implements.
- Well known Internet services include HTTP, FTP, DNS, finger, gopher, etc.
- User applications may also be built using the client-server paradigm.

# Client-Server: Session

- **Session**: the interaction between the server and one client.
- The service managed by a server may be accessed by multiple clients who desire the service, sometimes concurrently.
- Each client, when serviced by the server, engages in a **separate session** with the server, during which it conducts a dialog with the server until the client has obtained the service it required

# Client-Server





# Client-Server: Protocols

- A **protocol** is needed to specify the rules that must be observed by the client and the server during the conduction of a service.
- Such rules include specifications on matters such as
  - how the service is to be located
  - the sequence of interprocess communication
  - the representation and interpretation of data exchanged with each IPC
- On the Internet, such protocols are specified in the RFCs


# Client-Server: Locating a Service

- A mechanism must be available to allow a client process to locate a server for a given service.
- For example, in Internet services: address of the server process/hostname + protocol port number assigned.
- Well-known services such as FTP, HTTP, or telnet are assigned to a default port number reserved.
- Higher level of abstraction: service identified using a logical name registered with a registry.
- Logical name mapped to the physical location.
- If mapped at runtime (when a client process is run), then it is possible for the service's location to be dynamic, or moveable

# Client-Server: Implementing a service

- Implementations needs to adhere to the specification for the protocol, including how the dialogs of each session should proceed. The specification defines:
  - which side (client or server) should speak first
  - the syntax and semantic of each request and response
  - the action expected of each side upon receiving a particular request or response

# Client-Server: synchronization



# Session IPC examples

## Daytime service [RFC867]

- Client: Hello, <client address> here. May I have a timestamp please.
- Server: Here it is: (time stamp follows)

## HTTP [RFC2616]


- Client: Hello, <client address> here.
- Server: Okay. I am a web server and speaks protocol HTTP1.0
- Client: Great, please get me the web page index.html at the root of your document tree
- Server: Ok, here is the web page (content follows)

# Client-Server: Protocol Data Representation


- Part of the specification of a protocol is the synt ax and semantics of each request and response.
- Depends on the nature and the needs of the protocol.
- Representing data using text (character strings) is common, as it facilitates data marshalling and allows the data to be readable by human.
- Most well known Internet protocols are client- server, request-response, and text-based.

# Client-Server: software engineering

## **client-side software**



## **server-side software**



# Client server: advantages of separating the layers of logic

- It allows each module to be developed by people with special skills to focus on a module for which they have expertise.
  - UX software engineers may concentrate on developing the modules for the presentation logic
  - Other developers specialized in application logic and the service logic may focus on developing the other modules
- The separation allows modifications to be made to the logic at the presentation without requiring changes to be made at the lower layers.
- For example, the user interface can be changed from text-mode to graphical without requiring any change be made to the application logic or the service logic.



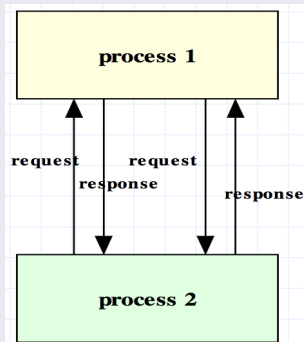
## Section 4 | Peer-to-peer

# The P2P System architecture

- An architecture where computer resources and services are direct exchanged between computer systems.
- Exchange of information, processing cycles, cache storage, and disk storage for files..
- Computers that have traditionally been used solely as clients communicate directly among themselves and can act as both clients and servers, assuming whatever role is most efficient for the network.

# P2P roles

- The participant processes play equal roles, with equivalent capabilities and responsibilities (hence the term “peer”).
- Each participant may issue a request to another participant and receive a response



# P2P usage and examples

- Client-server paradigm is an ideal model for a centralized network service,
- But Peer-to-peer paradigm is more appropriate for applications such as instant messaging, peer-to-peer file transfers, video conferencing, and collaborative work.
- It is also possible for an application to be based on both the client-server model and the peer-to-peer model
- Examples: Napster, GNutella, BitTorrent...


## Section 5 | Message System Paradigm

# Message System Paradigm

- It is an elaboration of the basic message-passing paradigm.
- A message system serves as an intermediary among separate, independent processes.
- The message system acts as a switch for messages, through which processes exchange messages asynchronously, in a decoupled manner.
- A sender deposits a message with the message system, which forwards it to a message queue associated with each receiver.
- Once a message is sent, the sender is free to move on to other tasks.
- Two subtypes:
  - Point-to-point
  - Publish/Subscribe


# Point to point

- A message system forwards a message from the sender to the receiver's message queue.
- Unlike the basic message passing model, the middleware provides:
  - A message depository , that allows the sending and the receiving to be decoupled.
  - Additional abstraction for asynchronous operations. Only using message-passing: a developer will have to make use of threads or child processes.
- Via the middleware, a sender deposits a message in the message queue of the receiving process.
- A receiving process extracts the messages from its message queue, and handles each one accordingly.



# Publish/Subscribe Message model

- Each message is associated with a specific **topic** or event.
- Applications interested in the occurrence of a specific event may subscribe to messages for that event.
- When the awaited event occurs, the process publishes a message announcing the event or topic.
- The middleware message system distributes the message to all its subscribers.
- The publish/subscribe message model offers a powerful abstraction for **multicasting** or group communication.





## Section 6 | Remote Procedure Call

# Remote Procedure Call

- A paradigm which allows distributed software to be programmed in a manner similar to conventional applications which run on a single processor.
- A remote procedure call involves two independent processes, which may reside on separate machines.
  - A process, A, wishing to make a request to another process, B, issues a procedure call to B, passing with the call a list of argument values.
  - As in the case of local procedure calls, a remote procedure call triggers a predefined action in a procedure provided by process B.
  - At the completion of the procedure, process B returns a value to process A.
- Tool: *rpcgen*

## Section 7 | Distributed Objects

# Distributed objects paradigms

- Applications access objects distributed over a network.
- Objects provide methods, through the invocation of which an application obtains access to services.

# Object-oriented paradigms

- Remote method invocation (RMI): object in a remote host
- Network services: using a service directory (example: Jini)
- Object request broker: CORBA
- Object spaces: publish/subscribe objects in a shared space

# References

- Distributed Systems. Marteen Van Steen and Andrew Tanenbaum (2017).
- Mei Ling-Liu. Distributed Computing Algorithms.
- Classroom notes by Guadalupe Ortiz